

# Фреймворк PyTorch для разработки искусственных нейронных сетей

## Урок 3. Dataset, Dataloader, BatchNorm, Dropout, Оптимизация

### Практическое задание

#### ▼ Попрактикуемся с тем, что изучили

Будем практиковаться на датасете: <https://www.kaggle.com/c/avito-demand-prediction>

Ваша задача:

1. Создать Dataset для загрузки данных (используем только числовые данные)
2. Обернуть его в Dataloader
3. Написать архитектуру сети, которая предсказывает число показов на основании числовых данных (вы всегда можете нагенерить дополнительных факторов). Сеть должна включать BatchNorm слои и Dropout (или НЕ включать, но нужно обосновать)
4. Учить будем на функцию потерь с катла (log RMSE) - нужно её реализовать
5. Сравните сходимость Adam, RMSProp и SGD, сделайте вывод по качеству работы модели

train-test разделение нужно сделать с помощью sklearn random\_state=13, test\_size = 0.25

Вопросы? в личку @Kinetikm

Выполнил **Соковнин ИЛ**

#### ▼ Как загрузить данные на Colaboratory

<https://www.youtube.com/watch?v=Ve5oW1qgbZg>

[https://colab.research.google.com/drive/1e3bbNmpQ1WremP4O5NiN0ggTV-WG5wL#scrollTo=1\\_lwnzf4Bhlu](https://colab.research.google.com/drive/1e3bbNmpQ1WremP4O5NiN0ggTV-WG5wL#scrollTo=1_lwnzf4Bhlu)

#### ▼ Шаг 0. Подготовка данных

Загружаем данные с Google Drive, затем распаковываем их в папку [/content/drive/my\\_drive/data](#)

```
!ls -la
```

```
total 16
drwxr-xr-x 1 root root 4096 Mar 23 14:22 .
drwxr-xr-x 1 root root 4096 Mar 30 08:53 ..
drwxr-xr-x 4 root root 4096 Mar 23 14:21 .config
drwxr-xr-x 1 root root 4096 Mar 23 14:22 sample_data
```

Создаём директорию для данных

```
!mkdir -p /content/drive/my\_drive/data
```

#### ▼ Использование Google Drive

Копирование больших файлов просходит значительно быстрее, если использовать Google Drive.

Подключение Google Drive к виртуальной машине:

```
from google.colab import drive
drive.mount('/content/gdrive', force_remount=True)
```

```
Mounted at /content/gdrive
```

```
!ls /content/gdrive/"My Drive" -la
```

```
total 2332713
drwx----- 2 root root      4096 Dec  9 2020 'Colab Notebooks'
-rw----- 1 root root 2388667849 Dec  4 16:43 features.csv.zip
drwx----- 2 root root      4096 Jan 21 08:43 nn
drwx----- 2 root root      4096 Mar 27 08:53 PyTorch
-rw----- 1 root root     15020 Oct 29 2020 task_1.drawio
-rw----- 1 root root      1767 Oct 29 2020 'Untitled Diagram.drawio'
-rw----- 1 root root       143 Apr  2 2021 Резюме.gdoc
```

```
!ls /content/gdrive/"My Drive"/PyTorch/lesson3/data -la
```

```
total 443208
-rw----- 1 root root   4144048 Mar 27 08:34 sample_submission.csv.zip
-rw----- 1 root root 116236318 Mar 28 08:46 test.csv.zip
-rw----- 1 root root 333463334 Mar 28 08:48 train.csv.zip
```

## Просматриваем подключенные диски

```
!df -h
```

Filesystem	Size	Used	Avail	Use%	Mounted on
overlay	108G	40G	69G	37%	/
tmpfs	64M	0	64M	0%	/dev
shm	5.8G	0	5.8G	0%	/dev/shm
/dev/root	2.0G	1.2G	817M	59%	/sbin/docker-init
tmpfs	6.4G	32K	6.4G	1%	/var/colab
/dev/sda1	81G	44G	38G	54%	/etc/hosts
tmpfs	6.4G	0	6.4G	0%	/proc/acpi
tmpfs	6.4G	0	6.4G	0%	/proc/scsi
tmpfs	6.4G	0	6.4G	0%	/sys/firmware
drive	15G	3.3G	12G	22%	/content/gdrive

```
!pwd
```

```
/content
```

## Просматриваем содержимое диска

```
!ls /content/gdrive/
```

```
MyDrive
```

```
!ls /content/gdrive/"My Drive"/PyTorch/lesson3/data -la
```

```
total 443208
-rw----- 1 root root   4144048 Mar 27 08:34 sample_submission.csv.zip
-rw----- 1 root root 116236318 Mar 28 08:46 test.csv.zip
-rw----- 1 root root 333463334 Mar 28 08:48 train.csv.zip
```

## Копирование данных с Google Drive на локальный диск виртуальной машины

```
# !cp /content/gdrive/'My Drive'/PyTorch/lesson3/data/*.* .
!cp /content/gdrive/'My Drive'/PyTorch/lesson3/data/train.csv.zip .
```

```
!ls -la
```

```
total 325676
drwxr-xr-x 1 root root      4096 Mar 30 08:55 .
drwxr-xr-x 1 root root      4096 Mar 30 08:53 ..
drwxr-xr-x 4 root root      4096 Mar 23 14:21 .config
drwxr-xr-x 3 root root      4096 Mar 30 08:54 drive
drwx----- 5 root root      4096 Mar 30 08:55 gdrive
drwxr-xr-x 1 root root      4096 Mar 23 14:22 sample_data
-rw----- 1 root root 333463334 Mar 30 08:55 train.csv.zip
```

```
!mv train.csv.zip /content/drive/my_drive/data/train.csv.zip
# !mv test.csv.zip /content/drive/my_drive/data/test.csv.zip
# !mv sample_submission.csv.zip /content/drive/my_drive/data/sample_submission.csv.zip
```

```
!ls -la
```

```
total 24
drwxr-xr-x 1 root root 4096 Mar 30 08:55 .
```

```
drwxr-xr-x 1 root root 4096 Mar 30 08:53 ..
drwxr-xr-x 4 root root 4096 Mar 23 14:21 .config
drwxr-xr-x 3 root root 4096 Mar 30 08:54 drive
drwx----- 5 root root 4096 Mar 30 08:55 gdrive
drwxr-xr-x 1 root root 4096 Mar 23 14:22 sample_data
```

## Распаковка

```
!ls -la /content/drive/my_drive/data
```

```
total 325660
drwxr-xr-x 2 root root      4096 Mar 30 08:55 .
drwxr-xr-x 3 root root      4096 Mar 30 08:54 ..
-rw----- 1 root root 333463334 Mar 30 08:55 train.csv.zip
```

```
!unzip '/content/drive/my_drive/data/train.csv.zip' -d '/content/drive/my_drive/data'
# !unzip '/content/drive/my_drive/data/test.csv.zip' -d '/content/drive/my_drive/data'
# !unzip '/content/drive/my_drive/data/sample_submission.csv.zip' -d '/content/drive/my_drive/data'
```

```
Archive:  /content/drive/my_drive/data/train.csv.zip
  inflating: /content/drive/my_drive/data/train.csv
```

## ▼ Import libs

```
import os, sys
import numpy as np
import pandas as pd
```

```
import torch
from torch import nn
from torch import optim
```

```
import torchvision
import torchvision.transforms as transforms
```

```
from tqdm import tqdm # means "progress" in Arabic (taqadum)
```

```
# import matplotlib.pyplot as plt ### воспользуемся для отображения изображения
```

```
!pwd
```

```
/content
```

## *File and column descriptions*

**train.csv** - Train data.

**item\_id** - Ad id.

**user\_id** - User id.

**region** - Ad region.

**city** - Ad city.

**parent\_category\_name** - Top level ad category as classified by Avito's ad model.

**category\_name** - Fine grain ad category as classified by Avito's ad model.

**param\_1** - Optional parameter from Avito's ad model.

**param\_2** - Optional parameter from Avito's ad model.

**param\_3** - Optional parameter from Avito's ad model.

**title** - Ad title.

**description** - Ad description.

**price** - Ad price.

**item\_seq\_number** - Ad sequential number for user.

**activation\_date** - Date ad was placed.

**user\_type** - User type.

**image** - Id code of image. Ties to a jpg file in train\_jpg. Not every ad has an image.

**image\_top\_1** - Avito's classification code for the image.

**deal\_probability** - The target variable.

This is the likelihood that an ad actually sold something. It's not possible to verify every transaction with certainty, so this column's value can be any float from zero to one.

## 1 Data loading

```
PATH_DATA = '/content/drive/my_drive/data/'
```

### 1.1 Train data

```
data_train = pd.read_csv(PATH_DATA + "train.csv")
# train = pd.read_csv(PATH_DATA + "train.csv")
# data_train = train.head(3000)
```

```
data_train.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1503424 entries, 0 to 1503423
Data columns (total 18 columns):
#   Column                Non-Null Count  Dtype
---  -
0   item_id               1503424 non-null object
1   user_id              1503424 non-null object
2   region               1503424 non-null object
3   city                 1503424 non-null object
4   parent_category_name 1503424 non-null object
5   category_name        1503424 non-null object
6   param_1              1441848 non-null object
7   param_2              848882 non-null object
8   param_3              640859 non-null object
9   title                1503424 non-null object
10  description           1387148 non-null object
11  price                 1418062 non-null float64
12  item_seq_number      1503424 non-null int64
13  activation_date       1503424 non-null object
14  user_type            1503424 non-null object
15  image                1390836 non-null object
16  image_top_1          1390836 non-null float64
17  deal_probability     1503424 non-null float64
dtypes: float64(3), int64(1), object(14)
memory usage: 206.5+ MB
```

```
data_train.head(3)
```

```
# data_train.dtypes
```

```
data_train['activation_date_int'] = pd.to_numeric(data_train['activation_date'].str.replace("-", "").astype(int)
data_train['activation_date'] = pd.to_datetime(data_train.activation_date)
data_train['day_of_month'] = data_train.activation_date.apply(lambda x: x.day)
data_train['day_of_week'] = data_train.activation_date.apply(lambda x: x.weekday())
data_train['month'] = data_train.activation_date.apply(lambda x: x.month)
# data_train['month'] = data_train['activation_date'].dt.month
```

```
data_train[['activation_date', 'activation_date_int', 'day_of_month', 'day_of_week', 'month']].head()
```

```
group = data_train.groupby('region')['deal_probability']  
# group.all()  
group.count()
```

```
region  
Алтайский край          41520  
Башкортостан            68291  
Белгородская область   28868  
Владимирская область   26741  
Волгоградская область  48998  
Воронежская область    44116  
Иркутская область      44030  
Калининградская область 32756  
Кемеровская область    44635  
Краснодарский край     141416  
Красноярский край      53442  
Нижегородская область  73643  
Новосибирская область  62486  
Омская область         42939  
Оренбургская область   29303  
Пермский край          62704  
Ростовская область     89995  
Самарская область      73407  
Саратовская область    49645  
Свердловская область   94475  
Ставропольский край    39187  
Татарстан              81284  
Тульская область       25733  
Тюменская область      35411  
Удмуртия               28537  
Ханты-Мансийский АО    28709  
Челябинская область    78339  
Ярославская область     32814  
Name: deal_probability, dtype: int64
```

```
group.describe()
```

```

from tqdm import notebook # IPython/Jupyter Notebook progressbar decorator for iterators.

cols = ['region', 'city', 'category_name', 'user_type']

for col in notebook.tqdm(cols):
    group = data_train.groupby(col)['deal_probability'] # deal_probability - вероятность сделки
    mean = group.mean()
    data_train[col + '__deal_probability_avg'] = data_train[col].map(mean) # convert the data to the desired format

data_train['_deal_probability'] = data_train['deal_probability']


data_train = data_train.drop([
    'item_id',
    'user_id',
    'region',
    'city',
    'parent_category_name',
    'category_name',
    'param_1',
    'param_2',
    'param_3',
    'title',
    'description',
    # 'price',
    # 'item_seq_number',
    'activation_date',
    'user_type',
    'image',
    # 'image_top_1',
    'deal_probability'
], axis=1)

cols

['region', 'city', 'category_name', 'user_type']

for col in data_train.columns:
    if data_train[col].isna().sum() > 0:
        data_train[col].fillna(data_train[col].median(), inplace=True)

# В итоге получаем
data_train.head(3)

```

## ▼ 1. Создать Dataset для загрузки данных (используем только числовые данные)

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
```

Чтобы изменить содержимое ячейки, дважды нажмите на нее (или выберите "Ввод")

```
# data_train = data_train.head(10000)

# разделим данные на train/test
# train-test разделение нужно сделать с помощью sklearn random_state=13, test_size = 0.25
X_train, X_test = train_test_split(data_train, test_size =0.25, random_state=13)

X_train.head(3)
```

```
idx=0
y_0 = X_train.iloc[idx, -1:]
print(y_0)
x_0 = X_train.iloc[idx, 0:-1]
print(x_0)
```

```
   _deal_probability    0.0
Name: 31723, dtype: float64
price                1.300000e+03
item_seq_number      6.000000e+00
image_top_1          1.056000e+03
activation_date_int   2.017032e+07
day_of_month         2.300000e+01
day_of_week          3.000000e+00
month                3.000000e+00
region__deal_probability_avg  1.470655e-01
city__deal_probability_avg    1.890690e-01
category_name__deal_probability_avg  1.784801e-01
user_type__deal_probability_avg  1.495572e-01
Name: 31723, dtype: float64
```

```
from torch.utils.data import Dataset
```

```
class AvitoDataset(Dataset):
    def __init__(self, init_dataset, normalize=False, transform=None):
        self._base_dataset = init_dataset
        self.normalize = normalize
        self.transform = transform

    def __len__(self):
        return len(self._base_dataset)

    def __getitem__(self, idx):
        """ Основной метод-точка входа для наборов данных PyTorch

        Аргументы:
            idx (int): индекс данных
        Возвращает:
            словарь признаков (x_data) и метки (y_target) данных
        """
        y = self._base_dataset.iloc[idx, -1:]
        # y = torch.tensor(y.values, dtype=torch.float)
        y = torch.FloatTensor(y.values)
        row = torch.FloatTensor(self._base_dataset.iloc[idx, 0:-1].values)

        return row, y
```

```
df_train, df_test = AvitoDataset(X_train), AvitoDataset(X_test)
```

## ▾ 2. Обернуть его в DataLoader

Чтобы изменить содержимое ячейки, дважды нажмите на нее (или выберите "Ввод")

```
BATCH_SIZE = 1024
train_loader = torch.utils.data.DataLoader(df_train,
                                           batch_size=BATCH_SIZE,
                                           shuffle=True,
                                           num_workers=1)

test_loader = torch.utils.data.DataLoader(df_test,
                                           batch_size=BATCH_SIZE,
                                           shuffle=True,
                                           num_workers=1)
```

## ▾ 3. Написать архитектуру сети, которая предсказывает число показов на основании числовых данных (вы всегда можете нагенерить дополнительных факторов). Сеть должна включать BatchNorm слои и Dropout (или НЕ включать, но нужно обосновать)

```
# import torch.nn.functional as F

# class FeedForwardNN(nn.Module):
#     def __init__(self, input_dim, hidden_dim):
#         ''' В init() - делаем необходимые настройки.

#         Аргументы:
#         input_dim - размерность вектора входных данных
#         hidden_dim - параметр для настройки скрытых слоёв
#         ...
#         super(FeedForwardNN, self).__init__()

#         self.bn1 = nn.BatchNorm1d(input_dim)
#         self.fc1 = nn.Linear(input_dim, 4*hidden_dim)
#         self.dp1 = nn.Dropout(0.25)

#         self.bn2 = nn.BatchNorm1d(4*hidden_dim)
#         self.fc2 = nn.Linear(4*hidden_dim, 3*hidden_dim)
#         self.dp2 = nn.Dropout(0.15)

#         self.bn3 = nn.BatchNorm1d(3*hidden_dim)
#         self.fc3 = nn.Linear(3*hidden_dim, 2*hidden_dim)
#         self.dp3 = nn.Dropout(0.15)

#         self.bn4 = nn.BatchNorm1d(2*hidden_dim)
#         self.fc4 = nn.Linear(2*hidden_dim, 1)

#     def forward(self, x):
#         x = self.bn1(x)
#         x = self.fc1(x)
#         # Задаём функцию потерь Leaky_relu
#         # x = F.leaky_relu(x, 0.05)
#         x = torch.tanh(x)
#         x = self.dp1(x)

#         x = self.bn2(x)
#         x = self.fc2(x)
#         # Задаём функцию потерь Leaky_relu
#         # x = F.leaky_relu(x, 0.05)
#         x = torch.tanh(x)
#         x = self.dp2(x)

#         x = self.bn3(x)
#         x = self.fc3(x)
#         # Задаём функцию потерь Leaky_relu
#         # x = F.leaky_relu(x, 0.05)
#         x = torch.tanh(x)
#         x = self.dp3(x)
```



```

#         x = self.bn4(x)
#         x = self.fc4(x)
#         x = torch.sigmoid(x)

#         return x

# def predict(self, x):
#     x = self.forward(x)
#     x = F.softmax(x)
#     return x

# net = FeedForwardNN(11, 4)
# print(net)

class FeedForwardNN(nn.Module):
    def __init__(self, input_dim, hidden_dim):
        super(FeedForwardNN, self).__init__()

        self.bn1 = nn.BatchNorm1d(input_dim)
        self.fc1 = nn.Linear(input_dim, 8*hidden_dim)
        self.dp1 = nn.Dropout(0.25)

        self.bn2 = nn.BatchNorm1d(8*hidden_dim)
        self.fc2 = nn.Linear(8*hidden_dim, 4*hidden_dim)
        self.dp2 = nn.Dropout(0.15)

        self.bn3 = nn.BatchNorm1d(4*hidden_dim)
        self.fc3 = nn.Linear(4*hidden_dim, 2*hidden_dim)
        self.dp3 = nn.Dropout(0.15)

        self.bn4 = nn.BatchNorm1d(2*hidden_dim)
        self.fc4 = nn.Linear(2*hidden_dim, 1)

    def forward(self, x):
        x = self.bn1(x)
        x = self.fc1(x)
        x = torch.tanh(x)
        x = self.dp1(x)

        x = self.bn2(x)
        x = self.fc2(x)
        x = torch.tanh(x)
        x = self.dp2(x)

        x = self.bn3(x)
        x = self.fc3(x)
        x = torch.tanh(x)
        x = self.dp3(x)

        x = self.bn4(x)
        x = self.fc4(x)
        x = torch.sigmoid(x)

        return x

```

## ▼ 4. Учить будем на функцию потерь с кагла (log RMSE) - нужно её реализовать

$$RMSLE = \sqrt{\frac{1}{n} \sum_{i=1}^n (\log y_i - \log \hat{y}_i)^2},$$

```

def log_rmse(y_pred, y_true):
    loss = torch.sqrt(torch.mean(
        (torch.log(y_pred+1)-torch.log(y_true+1))**2)
    )
    return loss

# def log_rmse(y_pred, y_true):
#     RMSLE = torch.sqrt(torch.mean(
#         (torch.log(y_pred)-torch.log(y_true))**2
#     ))

#     return RMSLE

# def log_rmse(y_true, y_pred):

```

```
# d_i = (K.log(y_pred) - K.log(y_true))
# loss1 = K.mean(K.square(d_i))
# loss2 = K.square(K.sum(K.flatten(d_i),axis=-1))/(K.cast_to_floatx(2) * K.square(K.cast_to_floatx(K.int_shape(K.flatten(d_i)))[0])
# loss = loss1 - loss2
# return loss
```

## 5. Сравните сходимость (Оптимизаторы) Adam, RMSProp и SGD, сделайте вывод по качеству работы модели

train-test разделение нужно сделать с помощью sklearn random\_state=13, test\_size = 0.25

```
epochs = 5
lr = 0.01
for batch, label in train_loader:
    print(batch.shape)
    break

    torch.Size([1024, 11])

def tr_loss(test_loader, model):

    batches = len(test_loader)
    loss = 0

    with torch.no_grad():
        for X, y in test_loader:
            y_pred = model(X)
            loss += log_rmse(y_pred, y)

    loss /= batches

    return loss

def train_proc(train_loader, test_loader, model, optimizer):

    for i, data in enumerate(train_loader):
        inputs, labels = data[0], data[1]

        y_pred = model(inputs)
        loss = log_rmse(y_pred, labels)

        # обнуляем градиент
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    return tr_loss(test_loader, model)
```

### ADAM

```
model = FeedForwardNN(11, 8)
print(model)

FeedForwardNN(
  (bn1): BatchNorm1d(11, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (fc1): Linear(in_features=11, out_features=64, bias=True)
  (dp1): Dropout(p=0.25, inplace=False)
  (bn2): BatchNorm1d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (fc2): Linear(in_features=64, out_features=32, bias=True)
  (dp2): Dropout(p=0.15, inplace=False)
  (bn3): BatchNorm1d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (fc3): Linear(in_features=32, out_features=16, bias=True)
  (dp3): Dropout(p=0.15, inplace=False)
  (bn4): BatchNorm1d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (fc4): Linear(in_features=16, out_features=1, bias=True)
)

optimizer = torch.optim.RMSprop(model.parameters(), lr=lr)
print(f'ADAM:\n')
for epoch in tqdm(range(epochs)):
    result = train_proc(train_loader,
```

```

        test_loader,
        model,
        optimizer
    )

print(f'Test_loss={result}')
```

ADAM:

20%	<div><div></div></div>		1/5	[11:24<45:36, 684.04s/it]	Test_loss=0.18010510504245758
40%	<div><div></div></div>		2/5	[22:52<34:20, 686.85s/it]	Test_loss=0.1796482801437378
60%	<div><div></div></div>		3/5	[34:31<23:04, 692.17s/it]	Test_loss=0.1796303391456604
80%	<div><div></div></div>		4/5	[46:23<11:40, 700.17s/it]	Test_loss=0.17977136373519897
100%	<div><div></div></div>		5/5	[58:27<00:00, 701.51s/it]	Test_loss=0.18032890558242798

## ▼ RMSProp

```

model = FeedForwardNN(11, 8)
optimizer = torch.optim.Adam(model.parameters(), lr=lr)
print(f'RMSprop:\n')
```

```

for epoch in tqdm(range(epochs)):
    result = train_proc(train_loader,
                        test_loader,
                        model,
                        optimizer
                    )

print(f'Test_loss={result}')
```

RMSprop:

20%	<div><div></div></div>		1/5	[11:52<47:28, 712.18s/it]	Test_loss=0.1798865646123886
40%	<div><div></div></div>		2/5	[23:56<35:58, 719.39s/it]	Test_loss=0.17945596575737
60%	<div><div></div></div>		3/5	[35:57<24:00, 720.25s/it]	Test_loss=0.17960643768310547
80%	<div><div></div></div>		4/5	[47:46<11:55, 715.83s/it]	Test_loss=0.17945215106010437
100%	<div><div></div></div>		5/5	[59:31<00:00, 714.30s/it]	Test_loss=0.17936037480831146

## ▼ SGD

```

model = FeedForwardNN(11, 8)
optimizer = torch.optim.SGD(model.parameters(), lr=lr)
print(f'SGD:\n')
```

```

for epoch in tqdm(range(epochs)):
    result = train_proc(train_loader,
                        test_loader,
                        model,
                        optimizer
                    )

print(f'Test_loss: {result}')
```

SGD:

20%	<div><div></div></div>		1/5	[11:40<46:42, 700.52s/it]	Test_loss: 0.19872649013996124
40%	<div><div></div></div>		2/5	[23:29<35:16, 705.33s/it]	Test_loss: 0.18643750250339508
60%	<div><div></div></div>		3/5	[35:04<23:21, 700.80s/it]	Test_loss: 0.1849292516708374
80%	<div><div></div></div>		4/5	[46:41<11:39, 699.26s/it]	Test_loss: 0.18434157967567444
100%	<div><div></div></div>		5/5	[58:12<00:00, 698.50s/it]	Test_loss: 0.1837015599012375

## ▼ Вывод:

Наилучшую сходимость и лучший результат наблюдаем при использовании оптимизатора RMSProp

```

# for epoch in tqdm(range(epochs)):
#     running_loss = 0.0
#     for i, data in enumerate(train_loader):
#         inputs, labels = data[0], data[1]

#         # обнуляем градиент
#         optimizer.zero_grad()
```

```

#         outputs = model(inputs)
#         loss = log_rmse(outputs, labels)
#         loss.backward()
#         optimizer.step()

#         # выводим статистику о процессе обучения
#         running_loss += loss.item()
#         if i % 300 == 0:     # печатаем каждые 300 mini-batches
#             print('[%d, %5d] loss: %.3f' %
#                   (epoch + 1, i + 1, running_loss / 300))
#             running_loss = 0.0

# print('Training is finished!')

```

Root Mean Squared Error (RMSE) Submissions are scored on the root mean squared error. RMSE is defined as:

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2},$$

where  $\hat{y}$  is the predicted value and  $y$  is the original value.

